

ML Engineering Tutorial Part 2

Tutor: Ralf Mayet (mayet@campus.tu-berlin.de)

Adaptive Systems Group, Humboldt University Berlin

Content / Goals

- Build simple regression models with Tensorflow.
- Build a simple machine learning pipeline (preprocessing, learning, evaluation)
- Build a classification multi-layer perceptron with Tensorflow.
- Extend it to use convolutional layers and observe the difference.

Sources

- [1] [Tensorflow and Keras Basic Regression](#)
- [2] [A line-by-line layman's guide to Linear Regression using TensorFlow](#) (adapted to TF2)
- [3] [Tensorflow Documentation: Loading MNIST](#)
- [4] [Basic classification: Classify images of clothing](#)
- [5] [Simple MNIST Convnet](#)

```
In [1]: # Importing packages we'll be using
import numpy as np
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import matplotlib.pyplot as plt
```

```
In [2]: # Helper functions for plotting:
def plotData(X, Y, predictions=None, title="Data Visualization"):
    plt.scatter(X, Y)
    if predictions is not None:
        plt.scatter(X, predictions)
    plt.title(title)
    plt.xlabel("X")
    plt.ylabel("Y")
    plt.show()

def plotImage(image):
    plt.imshow(image)
    plt.colorbar()
    plt.show()
```

Part 2A: Re-visiting regression


In [3]:

```
# Generating toy dataset
X = np.linspace(0, 2, 100)

# linear
y = 1.5 * X + np.random.randn(len(X)) * 0.3 + 0.5

# sinusoidal
# y = 1.5 * np.sin(X**2) + np.random.randn(len(X)) * 0.2 + 0.5

# Plot using our utility function
plotData(X, y)
```



In [4]:

```
# Making a model
# Ref activation functions: https://www.researchgate.net/profile/Junxi\_Feng

model = tf.keras.Sequential([
    layers.Dense(input_shape=[1,], units=1)
])

# model = tf.keras.Sequential([
#     layers.Dense(input_shape=[1,], units=1, activation="tanh"),
#     layers.Dense(units=4, activation="tanh"),
#     layers.Dense(units=1)
# ])

model.summary()
```


Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 1)	2

Total params: 2
Trainable params: 2
Non-trainable params: 0

```
In [5]: # Get some example predictions
predictions = model.predict(X)
plotData(X,y,predictions)


meanSquareError = ((y-predictions)**2).mean()
print("Mean Square Error: %.2f" % meanSquareError)
```



Mean Square Error: 1.91


```
In [6]: # Train the model
model.compile(optimizer=tf.optimizers.SGD(learning_rate=0.1), loss='mean_absolute_error')
history = model.fit(X,y, epochs=300, verbose=0)

# Plot the loss
plt.plot(history.history['loss'])
plt.show()
```



```
In [7]: # Evaluate the final predictions
predictions = model.predict(X)
plotData(X,y,predictions)

meanSquareError = ((y-predictions)**2).mean()
print("Mean Square Error: %.2f" % meanSquareError)
```



Mean Square Error: 1.48

```
In [8]: # We can inspect the layer weights after training and observe they
# match our toy data
model.layers[0].weights
```

```
Out[8]: [<tf.Variable 'dense/kernel:0' shape=(1, 1) dtype=float32, numpy=array([[1.
3991199]], dtype=float32)>,
<tf.Variable 'dense/bias:0' shape=(1,) dtype=float32, numpy=array([0.50000
083], dtype=float32)>]
```

Part 2B: Loading Images


```
In [9]: # Loading MNIST using built-in TF function
# Ref https://www.tensorflow.org/api_docs/python/tf/keras/datasets/mnist/load_data
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()

# Make sure images have shape (28, 28, 1)
x_train = np.expand_dims(x_train, -1)
x_test = np.expand_dims(x_test, -1)
```

```
In [10]: # Inspect dataset
print(x_train.shape)
print(y_train.shape)
print(x_test.shape)
print(y_test.shape)
```


```
(60000, 28, 28, 1)
(60000,)
(10000, 28, 28, 1)
(10000,)
```

```
In [11]: # Printing individual image and label
print(x_train[0, :, :, 0])
print(y_train[0])
```

3

```
In [13]: # Plot 25 examples
plt.figure(figsize=(8,8))
for i in range(25):
    plt.subplot(5,5,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(x_train[i,:,:,0])
    plt.xlabel(y_train[i])
plt.show()
```



Part 2C: Classification of Handwritten Digits

In [14]:

```
# We have our data in x_train, y_train (see above)

# Let's build a model for classification:
model_mlp = tf.keras.Sequential([
    layers.Flatten(input_shape=(28, 28, 1)),
    layers.Dense(128, activation='relu'),
    layers.Dense(10)
], name="mnist_mlp_model")

# Convolutional version
model_cnn = tf.keras.Sequential([
    keras.Input(shape=(28, 28, 1)),
    layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
    layers.MaxPooling2D(pool_size=(2, 2)),
    layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
    layers.MaxPooling2D(pool_size=(2, 2)),
    layers.Flatten(),
    layers.Dropout(0.5),
    layers.Dense(10)
], name="mnist_cnn_model")

model_mlp.summary()
model_cnn.summary()
```


Model: "mnist_mlp_model"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense_1 (Dense)	(None, 128)	100480
dense_2 (Dense)	(None, 10)	1290
Total params: 101,770		
Trainable params: 101,770		
Non-trainable params: 0		

Model: "mnist_cnn_model"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten_1 (Flatten)	(None, 1600)	0
dropout (Dropout)	(None, 1600)	0
dense_3 (Dense)	(None, 10)	16010
Total params: 34,826		
Trainable params: 34,826		
Non-trainable params: 0		

In [15]:

```
# Compile the models
model_mlp.compile(optimizer='adam',
                  loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                  metrics=['accuracy'])

model_cnn.compile(optimizer='adam',
                  loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                  metrics=['accuracy'])
```

In [16]:

```
# Evaluate Accuracy on test data
_, test_acc_mlp = model_mlp.evaluate(x_test, y_test, verbose=2)
_, test_acc_cnn = model_cnn.evaluate(x_test, y_test, verbose=2)
print('Test accuracy MLP:', test_acc_mlp)
print('Test accuracy CNN:', test_acc_cnn)
```

```
313/313 - 0s - loss: 155.0851 - accuracy: 0.1694
313/313 - 1s - loss: 45.1181 - accuracy: 0.0951
Test accuracy MLP: 0.16940000653266907
Test accuracy CNN: 0.09510000050067902
```

In [17]:

```
# Fit to data
history_mlp = model_mlp.fit(x_train, y_train, epochs=10)
history_cnn = model_cnn.fit(x_train, y_train, epochs=10)


# Plot the accuracy
plt.plot(history_mlp.history['accuracy'], label="MLP")
plt.plot(history_cnn.history['accuracy'], label="CNN")
plt.xlabel("epoch")
plt.ylabel("accuracy")
plt.legend()
plt.show()
```

```
Epoch 1/10
1875/1875 [=====] - 2s 1ms/step - loss: 2.4870 - accuracy: 0.8571
Epoch 2/10
1875/1875 [=====] - 2s 1ms/step - loss: 0.3628 - accuracy: 0.9115
Epoch 3/10
1875/1875 [=====] - 2s 1ms/step - loss: 0.2786 - accuracy: 0.9304
Epoch 4/10
1875/1875 [=====] - 2s 1ms/step - loss: 0.2497 - accuracy: 0.9376
Epoch 5/10
1875/1875 [=====] - 2s 1ms/step - loss: 0.2276 - accuracy: 0.9453
Epoch 6/10
1875/1875 [=====] - 3s 1ms/step - loss: 0.2147 - accuracy: 0.9472
Epoch 7/10
1875/1875 [=====] - 2s 1ms/step - loss: 0.2019 - accuracy: 0.9510
Epoch 8/10
1875/1875 [=====] - 2s 1ms/step - loss: 0.2010 - accuracy: 0.9526
Epoch 9/10
1875/1875 [=====] - 2s 1ms/step - loss: 0.1909 - accuracy: 0.9541
Epoch 10/10
1875/1875 [=====] - 2s 1ms/step - loss: 0.1870 - accuracy: 0.9559
Epoch 1/10
1875/1875 [=====] - 37s 20ms/step - loss: 0.5589 - accuracy: 0.9000
Epoch 2/10
1875/1875 [=====] - 32s 17ms/step - loss: 0.1368 - accuracy: 0.9589
Epoch 3/10
1875/1875 [=====] - 30s 16ms/step - loss: 0.1128 - accuracy: 0.9665
Epoch 4/10
1875/1875 [=====] - 32s 17ms/step - loss: 0.0996 - accuracy: 0.9700
Epoch 5/10
1875/1875 [=====] - 32s 17ms/step - loss: 0.0907 - accuracy: 0.9728
Epoch 6/10
1875/1875 [=====] - 34s 18ms/step - loss: 0.0816 - accuracy: 0.9758
Epoch 7/10
1875/1875 [=====] - 34s 18ms/step - loss: 0.0772 -
```

```

accuracy: 0.9770
Epoch 8/10
1875/1875 [=====] - 33s 18ms/step - loss: 0.0752 -
accuracy: 0.9775
Epoch 9/10
1875/1875 [=====] - 31s 17ms/step - loss: 0.0710 -
accuracy: 0.9785
Epoch 10/10
1875/1875 [=====] - 32s 17ms/step - loss: 0.0667 -
accuracy: 0.9797

```



In [18]:

```

# Evaluate Accuracy on test data
_, test_acc_mlp = model_mlp.evaluate(x_test, y_test, verbose=2)
_, test_acc_cnn = model_cnn.evaluate(x_test, y_test, verbose=2)
print('Test accuracy MLP:', test_acc_mlp)
print('Test accuracy CNN:', test_acc_cnn)

```

```

313/313 - 0s - loss: 0.2574 - accuracy: 0.9535
313/313 - 1s - loss: 0.0451 - accuracy: 0.9859
Test accuracy MLP: 0.953499972820282
Test accuracy CNN: 0.9858999848365784

```

Part 2D: Evaluating the classification output

In [19]:

```

# Get some predictions
# Attach a softmax layer to convert the logits to probabilities, which are
probability_model = tf.keras.Sequential([model_cnn, tf.keras.layers.Softmax])
predictions = probability_model.predict(x_test)

print(predictions.shape)
print(predictions[0])

```




```




(10000, 10)
[1.11056151e-08 1.24207344e-14 4.61679619e-08 3.08107388e-08
 1.85735148e-16 1.80415318e-12 3.65252618e-20 1.00000000e+00
 9.83751830e-11 1.84875182e-09]

```

In [20]:

```
# Plot classification results
for i in range(3):
    plotImage(x_test[i,:,:,0])
    plt.xticks(range(10))
    plt.bar(range(10), predictions[i])
    plt.show()
```





Part 2E: Loading images from disk

Demo on Desktop